reveals

SevenBridges

Applied Bioinformatics

Applied Bioinformatics

The Reference Genome

Mladen Lazarevic
Milan Kovacevic
milan.kovacevic@sbgenomics.com

© 2016 Seven Bridges

Agenda

- PCR and Sequencing - recapitulation
- Human Genome Project
- File formats related to the reference genome
- Example tasks

© 2016 Seven Bridges

PCR - Polymerase chain reaction

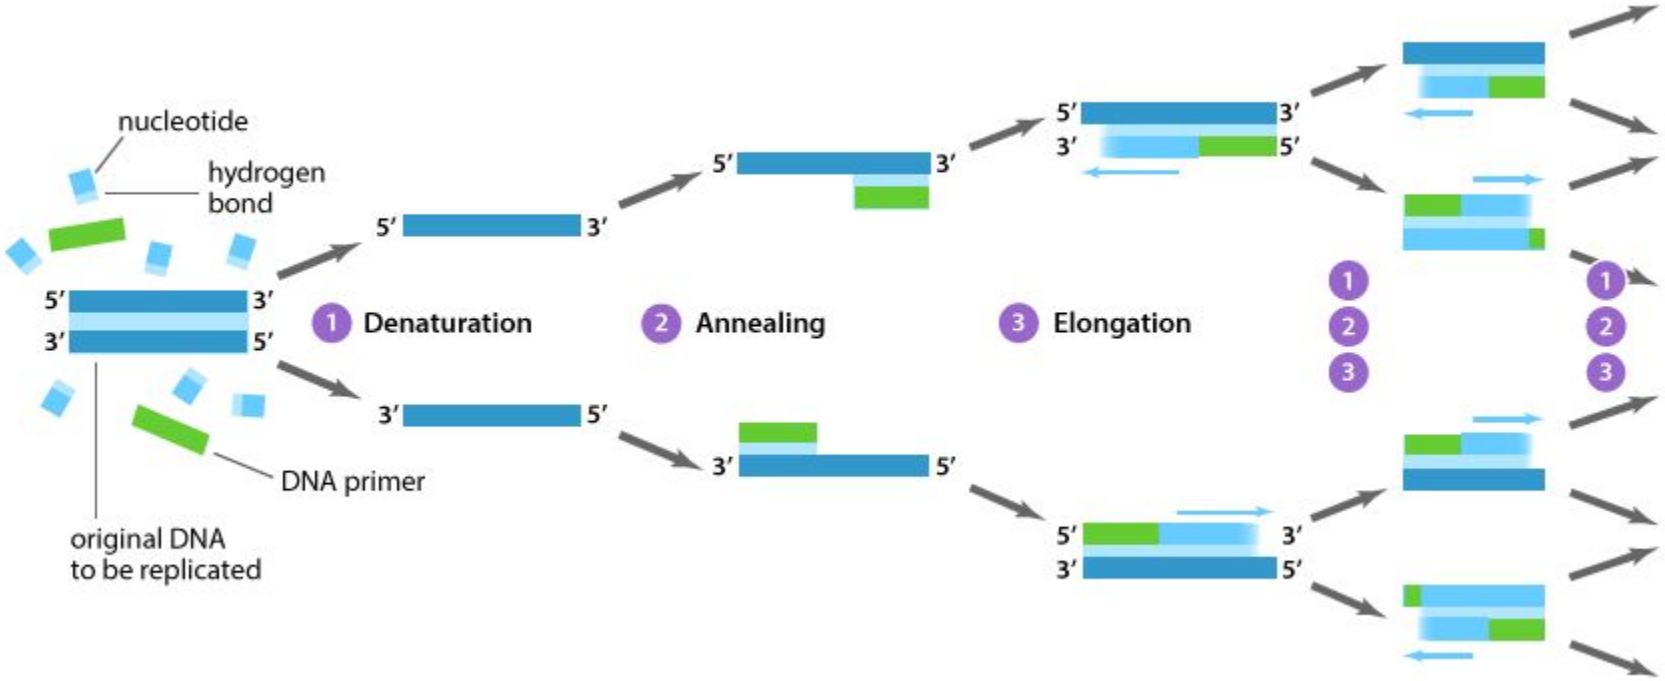

© 2016 Seven Bridges

PCR - Polymerase chain reaction

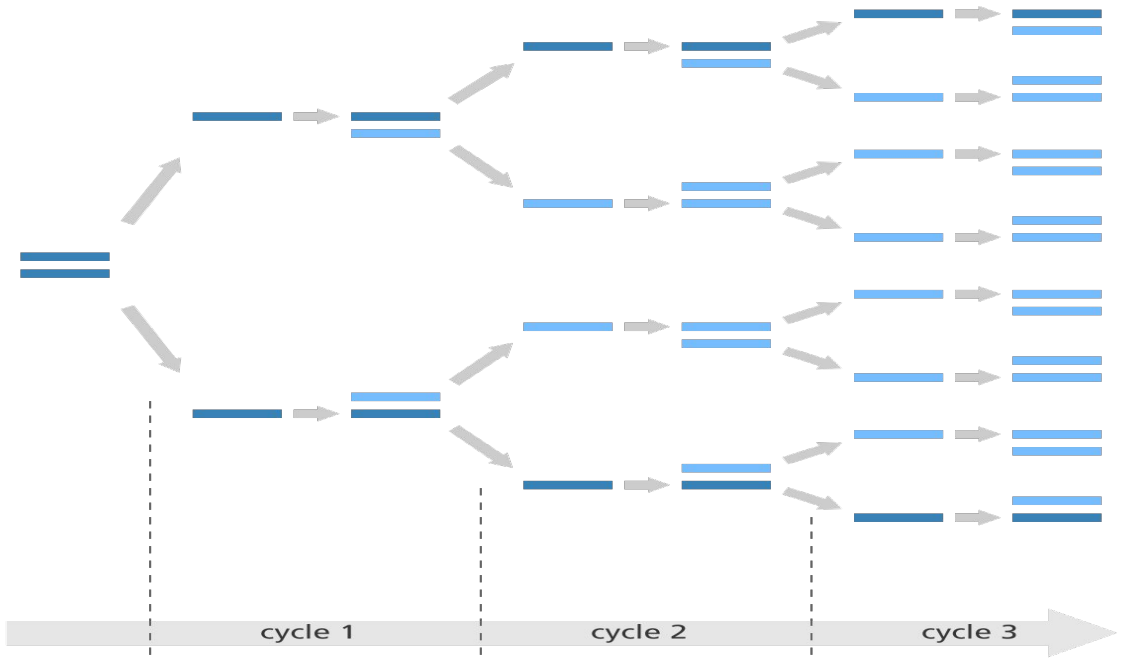

© 2016 Seven Bridges

PCR - Polymerase chain reaction

© 2016 Seven Bridges

PCR - Error

- 1 in 10k error rate

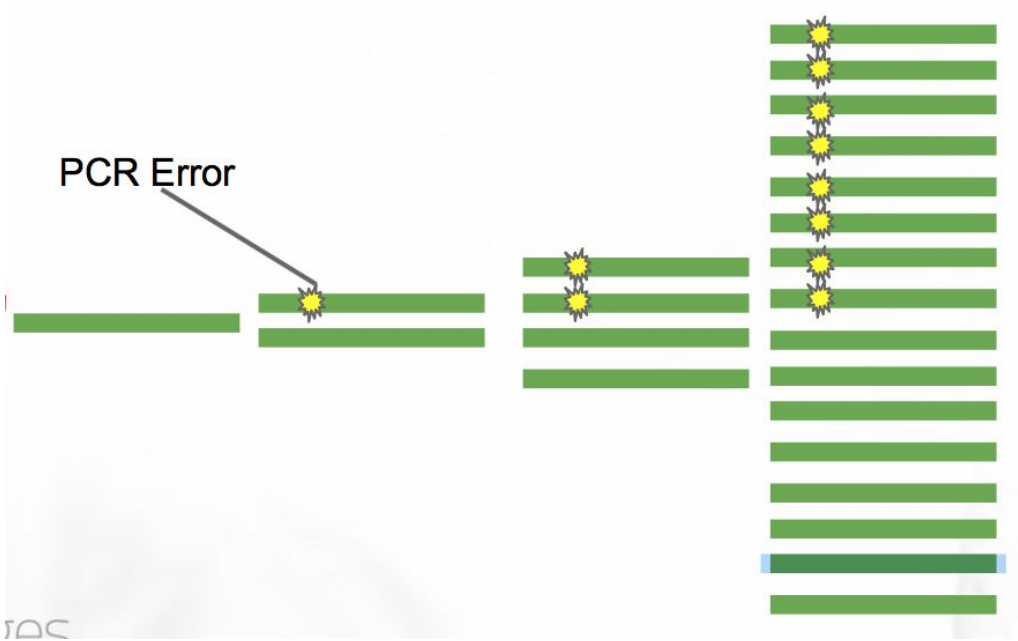

© 2016 Seven Bridges

Genome sequencing - a reminder

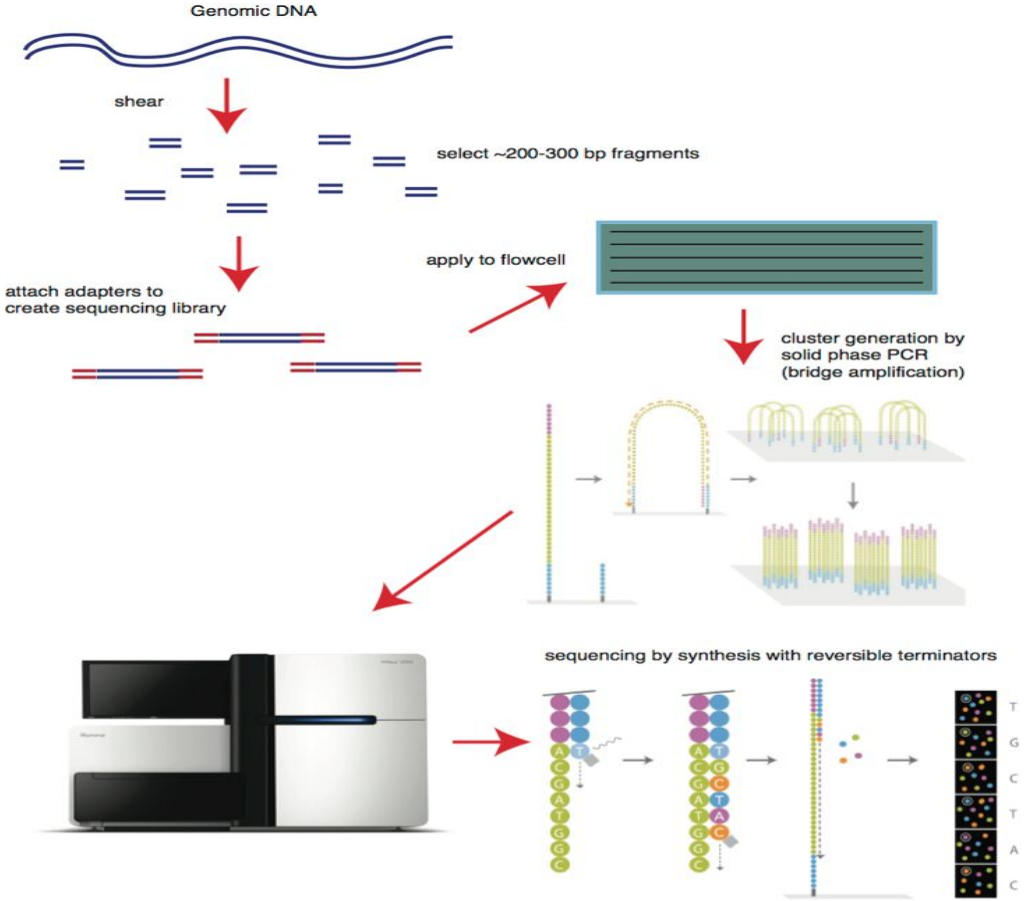

© 2016 Seven Bridges

Genome sequencing - recapitulation

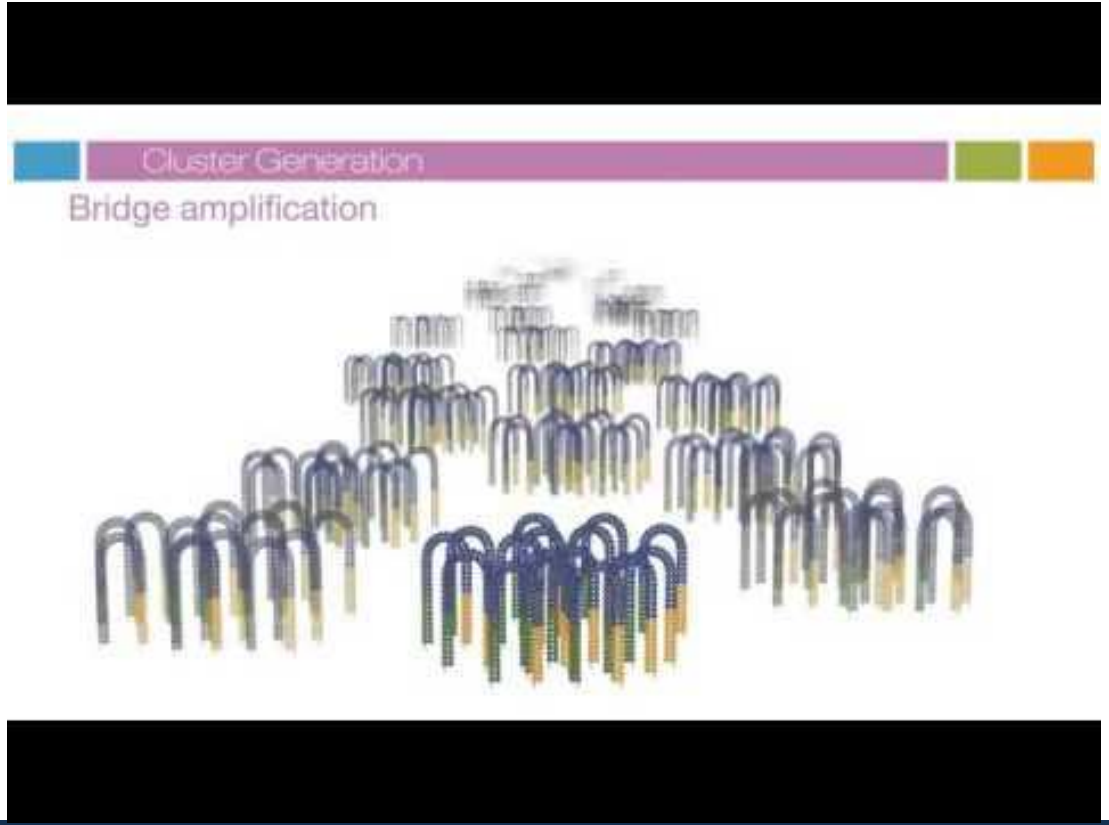

© 2016 Seven Bridges

Genome sequencing - recapitulation

Figure 4. Paired-End Sequencing and Alignment

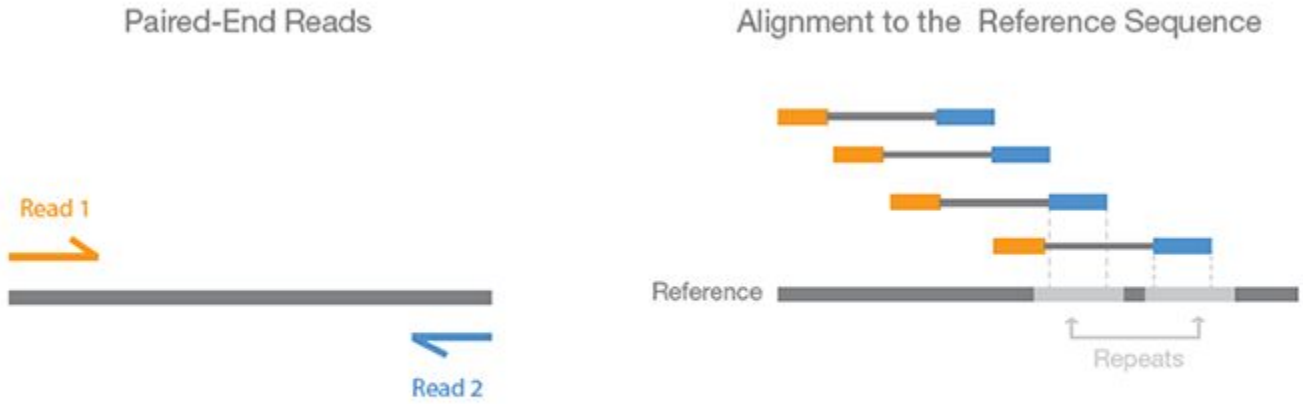

Paired-end sequencing enables both ends of the DNA fragment to be sequenced. Because the distance between each paired read is known, alignment algorithms can use this information to map the reads over repetitive regions more precisely. This results in much better alignment of the reads, especially across difficult-to-sequence, repetitive regions of the genome.

© 2016 Seven Bridges

DNA Sequencing - Reminder

- We got a FASTQ file with the “reads” - little pieces of the genome

© 2016 Seven Bridges

What to do with the sequencing reads?

- How do we reconstruct the genome that went into the “shredder”?
- We could try “assembly” - connecting the reads into longer sequences

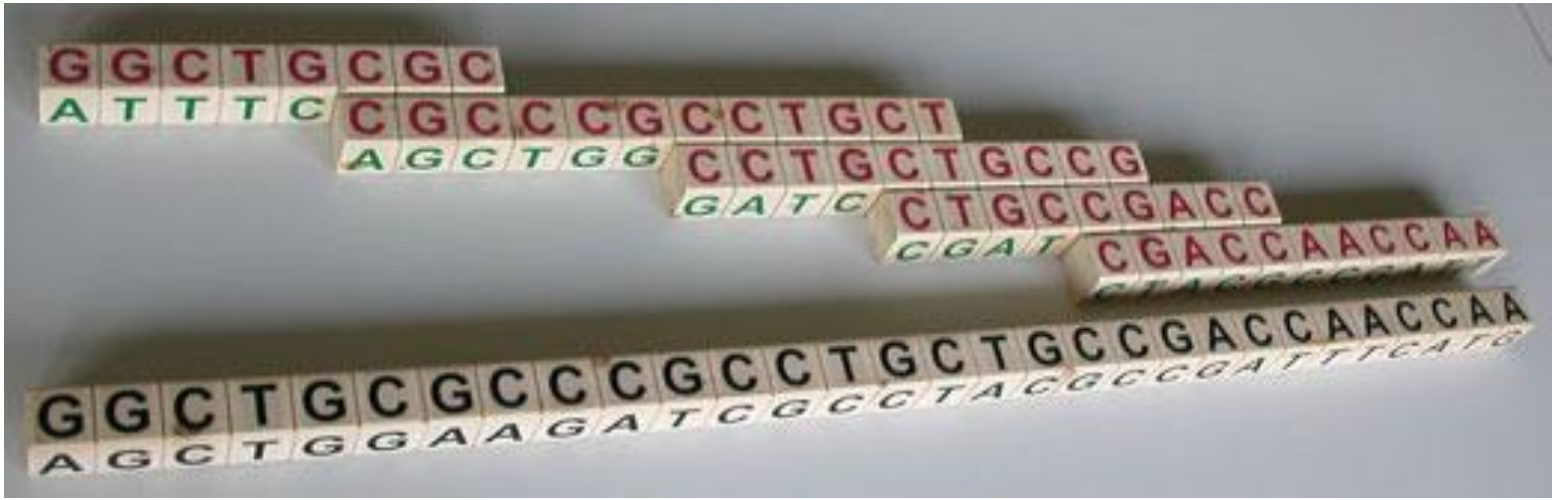

© 2016 Seven Bridges

Genome Assembly

- Greedy algorithm (suboptimal solution):
  1. Calculate pairwise alignments of all fragments.
  2. Choose two fragments with the largest overlap.
  3. Merge chosen fragments.
  4. Repeat step 2 and 3 until only one fragment is left.
- Even the more practical solutions have problems:
  - High computational cost
  - High memory consumption (100s of GB or RAM)

© 2016 Seven Bridges

Genome Assembly

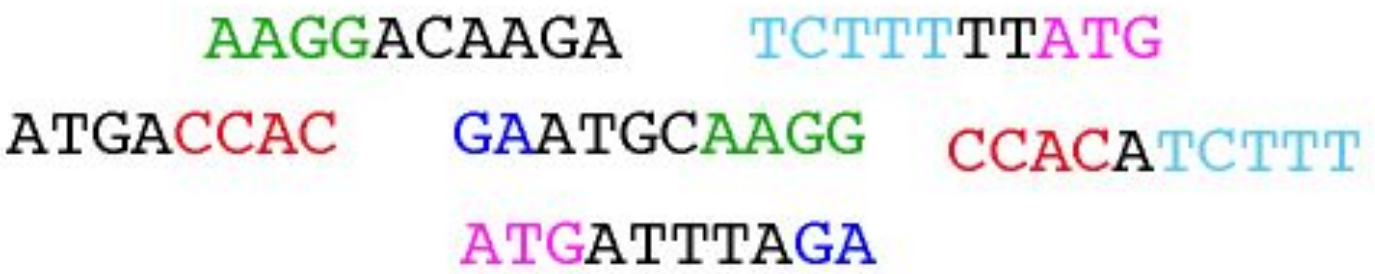

© 2016 Seven Bridges

What to do with reads? (an alternative)

- We do have an alternative approach to recover that genome (the genome that went into the “shredder”)
- This way is faster and more practical than assembly
- It is based around a *Reference genome*, *Alignment*, and *Variant Calling*
- But first we will need to define these terms and procedures

© 2016 Seven Bridges

The reference genome

- A reference genome is representative example of a species' genome
- It is often built from genomes of multiple individuals
- Every individual differs from this genome in some places
- We can think of *genomic variation* as differences from the reference (genome)
- Reference genomes provide a coordinate system for communicating genomic data
- And, they make analyses easier!

© 2016 Seven Bridges

What do with the sequencing reads? (an alternative)

© 2016 Seven Bridges

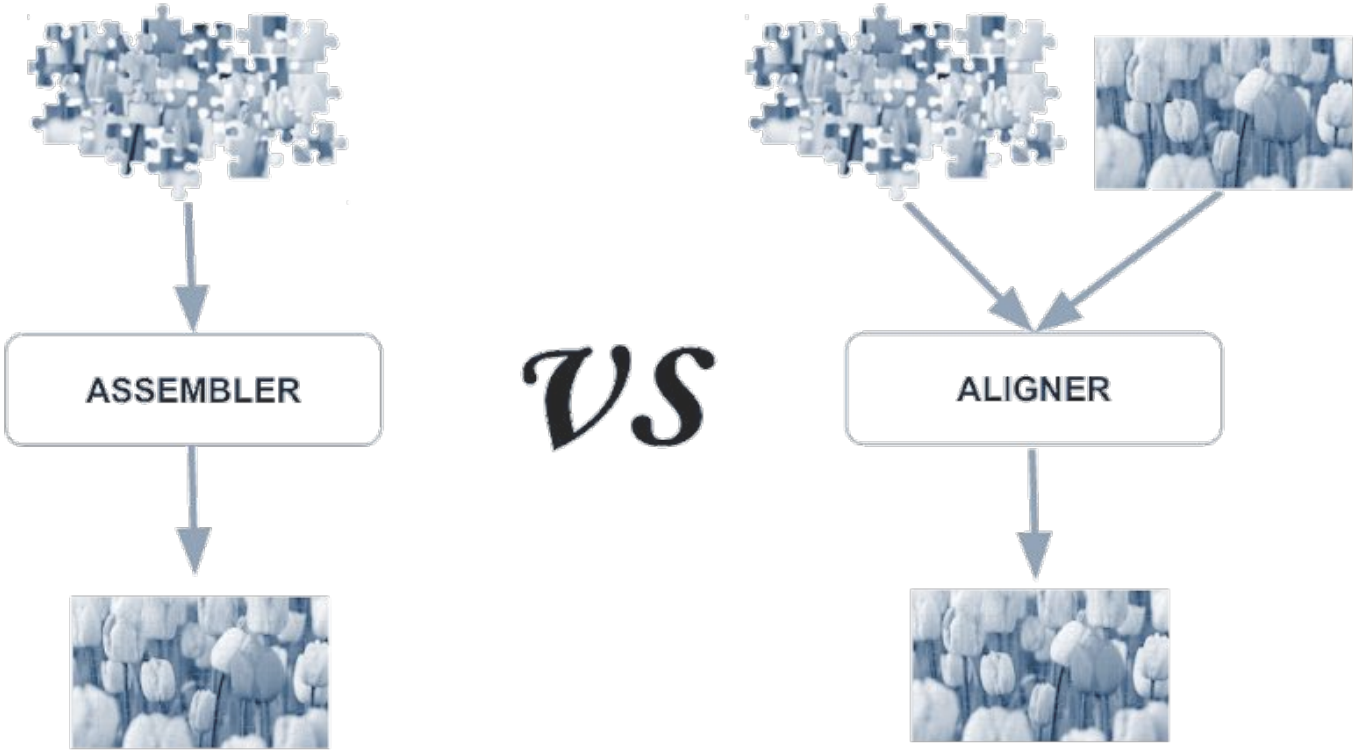

© 2016 Seven Bridges

Human Genome Project

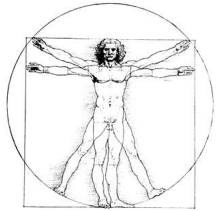

- International scientific initiative
  to create a reference human genome
- Active in years 1990 to 2003, at a cost over $3B
- 70% of the reference came from a single male donor
- In parallel, Celera Corporation launched a privately funded project,
  with the same goal, but had the intention to patent the sequence
- This caused a race, leading up the release of the first draft
  of the public version of the human genome on July 7, 2000,
  by the UCSC Genome Bioinformatics Group

© 2016 Seven Bridges

Human Genome Project

- 130 Books
- Font size: 4
- 43k letters per page

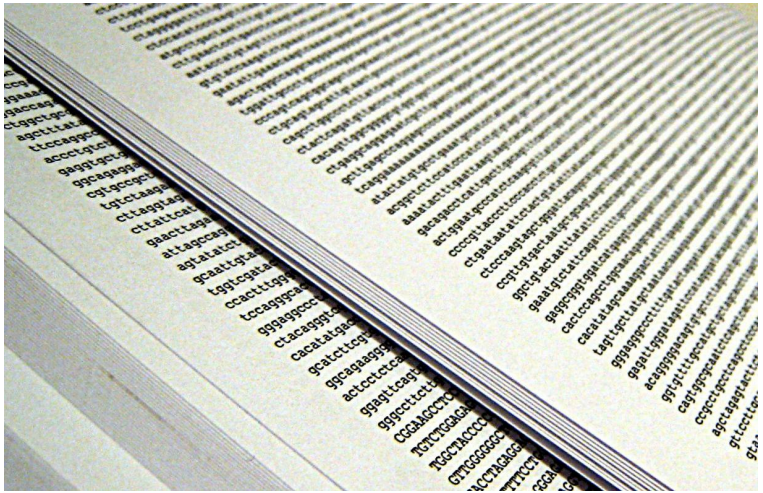

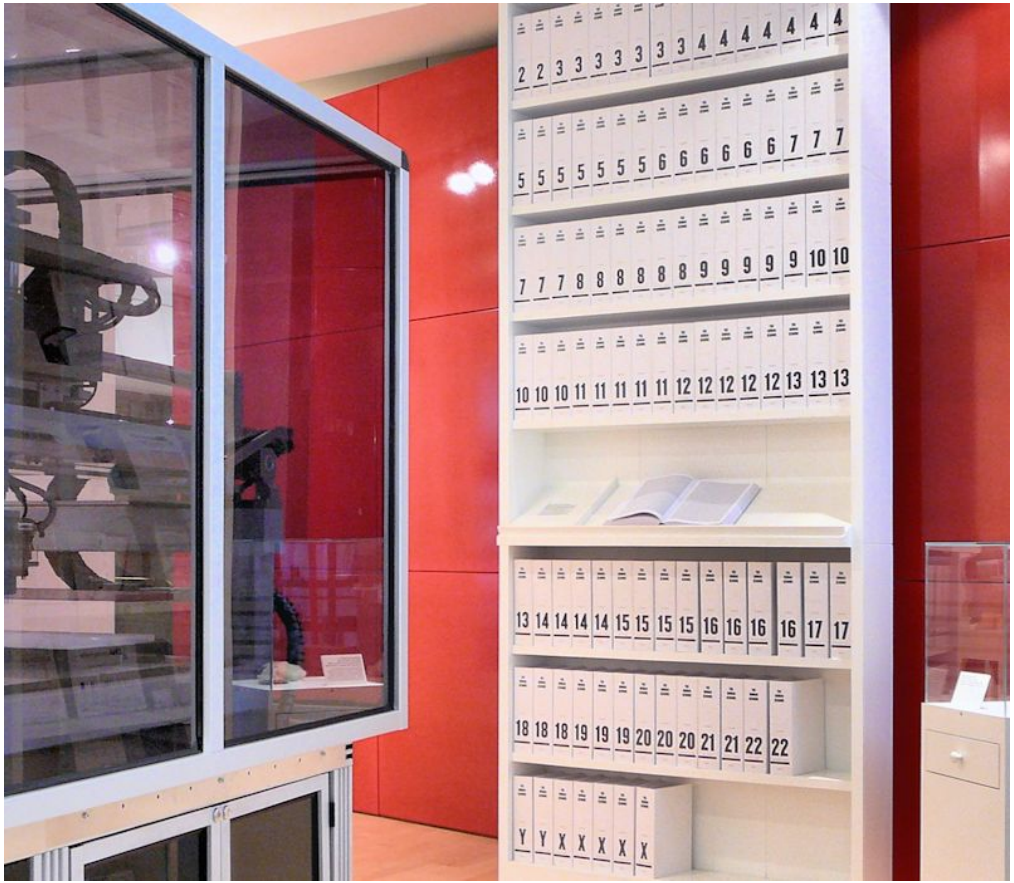

© 2016 Seven Bridges

Human genome stats

- The human genome has 23 pairs of chromosomes:
  - Males have 22 autosomes, one X and one Y chromosome
  - Females have 22 autosomes and two X chromosomes
- There is also a separate Mitochondrial DNA contig
- Total length is ~3 Gigabases (3B basepairs)
- About 20 000 genes covering about 30 Megabases
  - 2% is the coding region

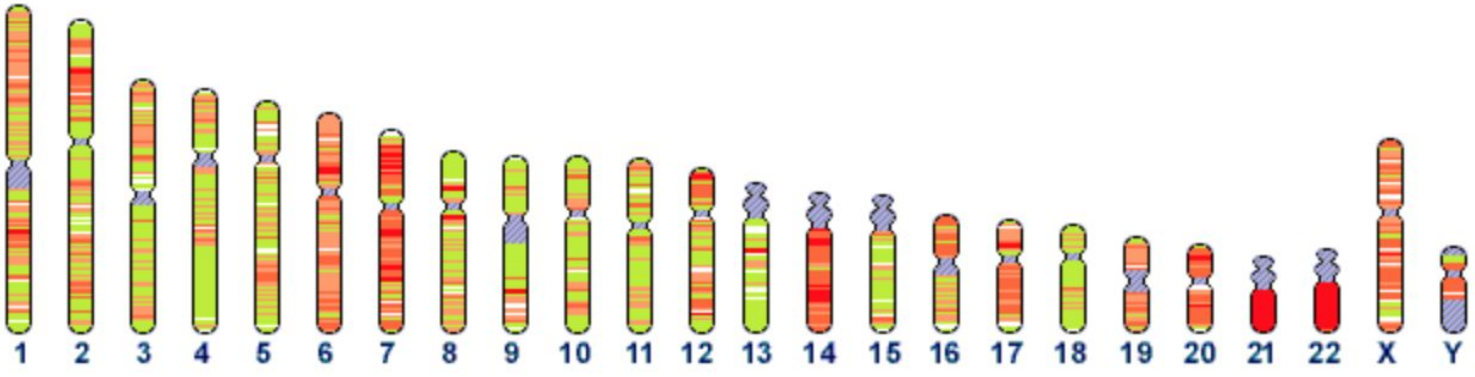

© 2016 Seven Bridges

Human exome stats

- Exome - part of the human genome coding for proteins/mRNA
- Covers only 2% of genome - much cheaper to sequence

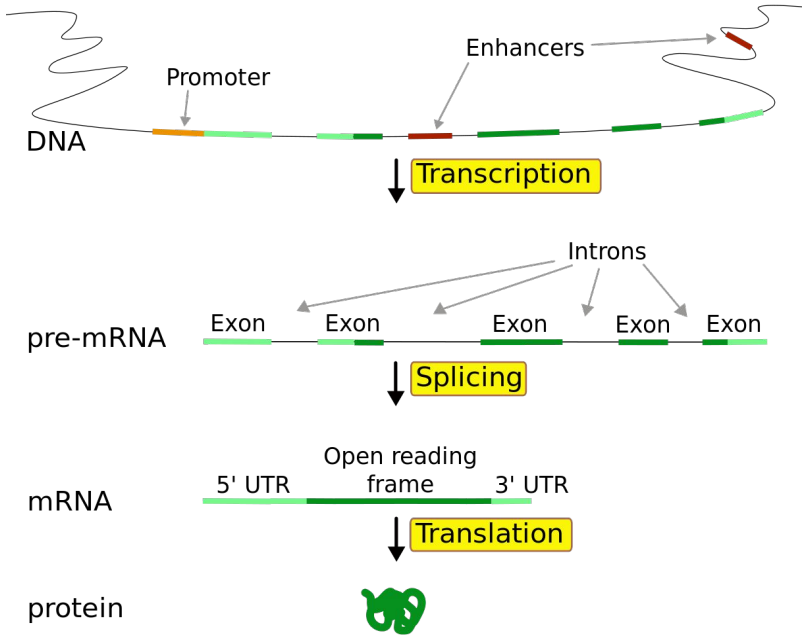

© 2016 Seven Bridges

Human genome HG38

- Current version of the chromosomes is *HG38* (though 37 is still used sometimes)
- Released December 24th 2013
- Additional sequences are added to the genome:
  - Unplaced sequences (some genomes contain them, somewhere)
  - Unlocalized sequences (chromosome known, but coordinates are not)
  - Alternate sequences (some genomes contain them, instead of some parts)
  - Human Herpesvirus 4 type 1
- Patches are commonly added
- http://www.ncbi.nlm.nih.gov/projects/genome/assembly/grc/human/

© 2016 Seven Bridges

FASTA file format

- A simple format for storing reference files
- Two “types” of lines:
  - Description lines, start with ‘>’, contain sequence name and optional description
  - Sequence lines follow a description line and contain the actual nucleotide sequence
- Sequences are represented by a single description line, followed by one or more sequence lines (usually 70 bases or less in a line)

```
>chr 1
ACTGCCCTCCGTACCTACTGCCCTCCGTACCTACTGCCCTCCGTACCTACTGCCCTCCGTACCT
ACTGCCCTCCGTACCTACTGCCCTCCGTACCTACTGCCCTCCGTACCTACTGCCCTCCGTACCT
ACTGCCCTCCGTACCTACTGCCCTCCGTACCTACTGCCCTCCGTACCTACTGCCCTCCGTACCT
ACTGCCCTCCGTACCTACTGCCCTCCGTACCTACTGCCCTCCG
```

© 2016 Seven Bridges

FASTA File IUPAC Codes (and contig names)

- A, C, T, G, U - Nucleotides (Adenine, Cytosine, Guanine, Thymine, Uracil)
- Ambiguous bases:
  - R - A or G
  - Y - C, T, or U
  - N - Any base
  - https://en.wikipedia.org/wiki/FASTA_format

- Often two versions of the Human Reference Genome are found
  - Chromosomes labeled 1, 2, 3… (Human Genome Consortium style)
  - Chromosomes labeled chr1, chr2, chr3… (UCSC style)

© 2016 Seven Bridges

FASTA index (fai)

- Some tools build or require indices of the fasta file
  - Needs to be in the same folder, and have the same name as the fasta + .fai extension
- Fai file structure:

1. The name of the sequence
2. The length of the sequence
3. The byte offset of the first base in the file
4. The number of bases in each fasta line
5. The number of bytes in each fasta line

Example:

```
1    249250621   52          60   61
2    243199373   253404903   60   61
3    198022430   500657651   60   61
```

© 2016 Seven Bridges

Pysam - Python fasta interface

- Pysam - a Python toolkit for working with genomic files
- pysam.Fastafile:
  - Create a fasta file parser: `fasta = pysam.Fastafile(path_to_file)`
  - Get the sequence names in the file: `fasta.references`
  - Get the lengths of the sequences: `fasta.lengths`
  - Retrieve a (part of) sequence: `fasta.fetch(sequence_name, [start], [stop])`
- Fasta coordinates in pysam are zero-based
  - Not true for all file types in that pysam supports
- Other fasta interfaces for Python exist
  - pyfasta - works only on fasta files
  - biopython - a much larger toolkit that complement sequences, etc.
  - We are describing pysam, as it covers all the file types covered in the course

© 2016 Seven Bridges

Exercise: FASTA file format (10 minutes)

- An example FASTA file is found under */sbgenomics/project-files/example_human_reference.fasta*
- View the contents of the file
  - You can use “!head filename.ext” in the Notebook to invoke linux head
- Create a pysam Fastafile parser
- Get and print sequence names
  - How many sequences are there?
- Fetch the entire sequence
  - How long is it?
  - Print the first 100 bases

© 2016 Seven Bridges

Exercise: FASTA file format (20 minutes)

- Full hg38 FASTA file is located under in */sbgenomics/project-files/Homo_sapiens_assembly38.fasta*
- Create a pysam Fastafile parser
- Get sequence names for all contigs
  - How many contigs are there?
  - Read the names of the contigs
  - How long is chromosome 5?
- Fetch section chromosome 17:43044295-43125370
  - What is the “GC content” on this region? (percent of G and C bases)
  - What is the most common 3-mer?
- Fetch base at chromosome 1:248755121
  - What is the base?
- Fetch region at chromosome 1:50000-50100
  - What is the base composition?

© 2016 Seven Bridges

FASTQ file format

- Most common format for storing sequencing reads
- It's spread across four lines. The four lines are:
  - “@” followed by a read name
  - Nucleotide sequence
  - “+” , possibly followed by some info, but ignored by most of the tools
  - Quality sequence

- Example:

```
@ERR294379.100739024 HS24_09441:8:2203:17450:94030#42/1
AGGGAGTCCACAGCACAGTCCAGACTCCCACCAGTTCTGACGAAATGATGAGAGCTCAGA
+
BDDEEF?FGFFFHGFFHHGHGGHCH@GHHHGFAHEGFEHGEFGHCCGGGFEGFGFFDFFHB
```

© 2016 Seven Bridges

Annotations file formats

- Annotations are stored in few (similar) file types
  - **BED**
  - GTF
  - GFF
  - ...
- **BED file format** (tab-separated):

CHROM START END NAME score ticks blocks

- Chrom, start and end are required
- Start is 0-based
- End is open interval (not included)

© 2016 Seven Bridges