

SevenBridges

Applied Bioinformatics

Applied Bioinformatics

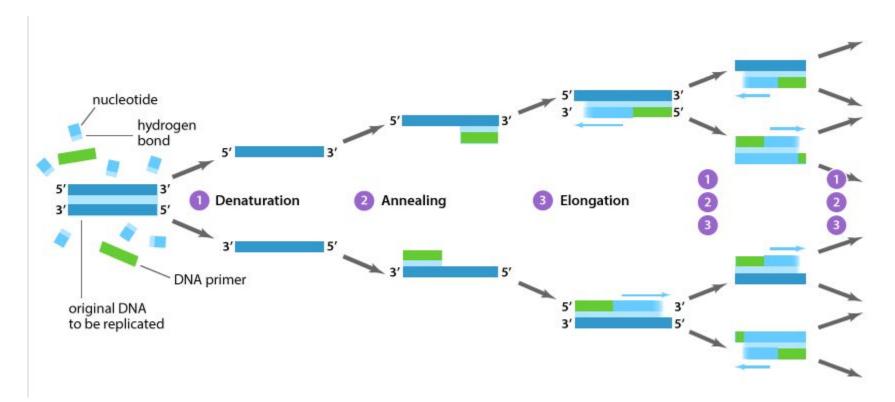
The Reference Genome

Mladen Lazarevic
Milan Kovacevic
milan.kovacevic@sbgenomics.com

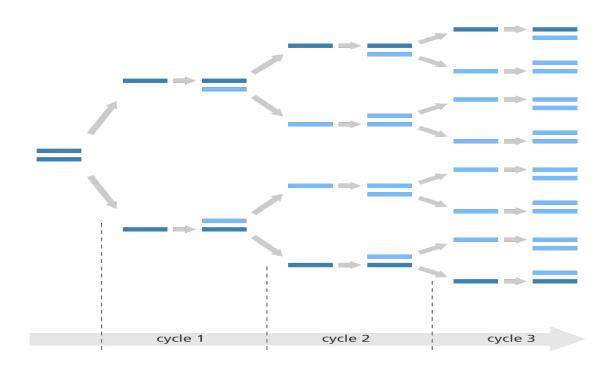
Agenda

- PCR and Sequencing recapitulation
- Human Genome Project
- File formats related to the reference genome
- Example tasks

PCR - Polymerase chain reaction



PCR - Polymerase chain reaction

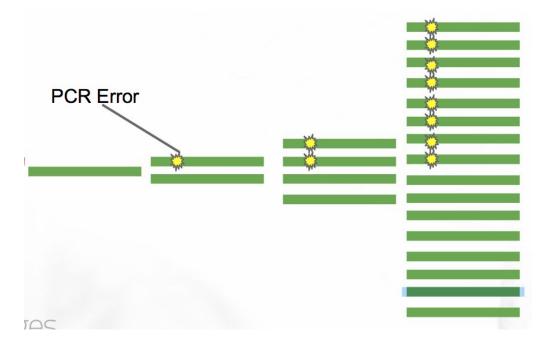


PCR - Polymerase chain reaction

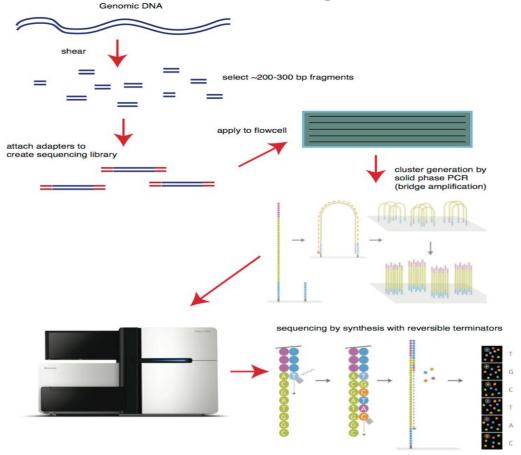


PCR - Error

• 1 in 10k error rate



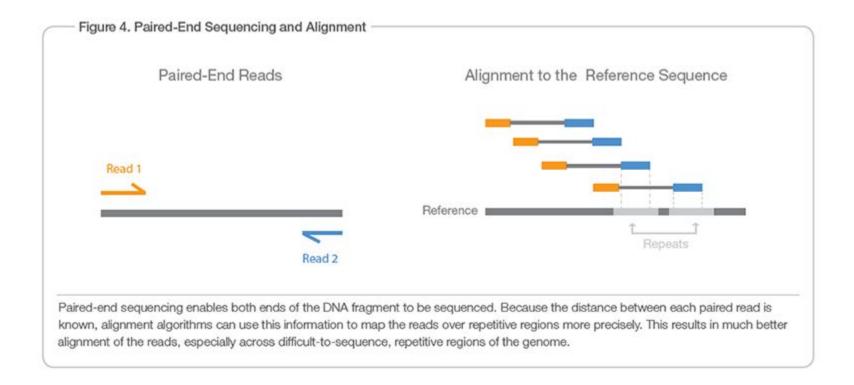
Genome sequencing - a reminder



Genome sequencing - recapitulation



Genome sequencing - recapitulation



DNA Sequencing - Reminder

■ We got a FASTQ file with the "reads" - little pieces of the genome



What to do with the sequencing reads?

- How do we reconstruct the genome that went into the "shredder"?
- We could try "assembly" connecting the reads into longer sequences



Genome Assembly

- Greedy algorithm (suboptimal solution):
 - 1. Calculate pairwise alignments of all fragments.
 - 2. Choose two fragments with the largest overlap.
 - 3. Merge chosen fragments.
 - 4. Repeat step 2 and 3 until only one fragment is left.
- Even the more practical solutions have problems:
 - High computational cost
 - High memory consumption (100s of GB or RAM)

Genome Assembly

AAGGACAAGA

TCTTTTTATG

ATGACCAC

GAATGCAAGG

CCACATCTTT

ATGATTTAGA

What to do with reads? (an alternative)

- We do have an alternative approach to recover that genome (the genome that went into the "shredder")
- This way is faster and more practical than assembly
- It is based around a Reference genome, Alignment, and Variant Calling
- But first we will need to define these terms and procedures

The reference genome

- A reference genome is representative example of a species' genome
- It is often built from genomes of multiple individuals
- Every individual differs from this genome in some places
- We can think of genomic variation as differences from the reference (genome)
- Reference genomes provide a coordinate system for communicating genomic data
- And, they make analyses easier!

What do with the sequencing reads? (an alternative)

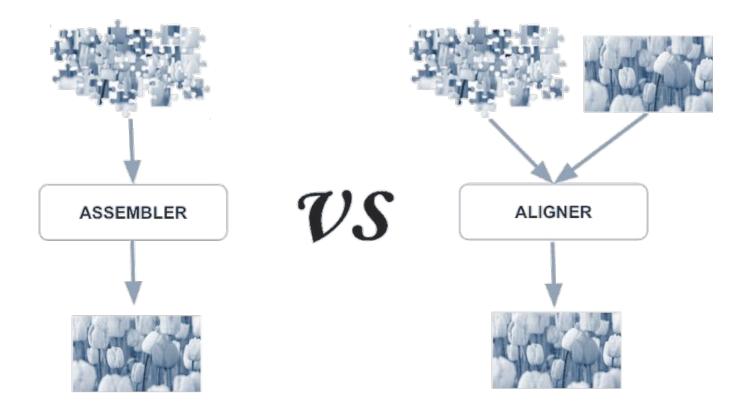
AAGGACAAGA TCTTTTTATG

ATGACCAC

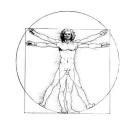
GAATGCAAGG

CCACATCTTT

ATGATTTAGA



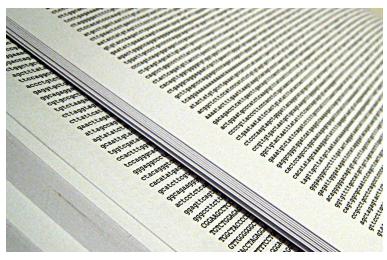
Human Genome Project



- International scientific initiative
 to create a reference human genome
- Active in years 1990 to 2003, at a cost over \$3B
- 70% of the reference came from a single male donor
- In parallel, Celera Corporation launched a privately funded project, with the same goal, but had the intention to patent the sequence
- This caused a race, leading up the release of the first draft of the public version of the human genome on July 7, 2000, by the UCSC Genome Bioinformatics Group

Human Genome Project

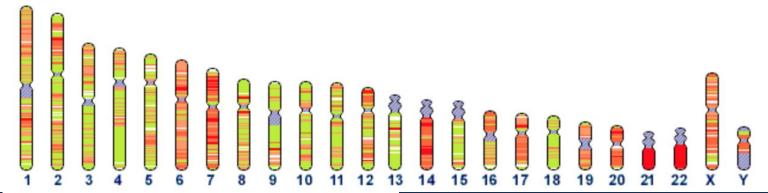
- 130 Books
- Font size: 4
- 43k letters per page





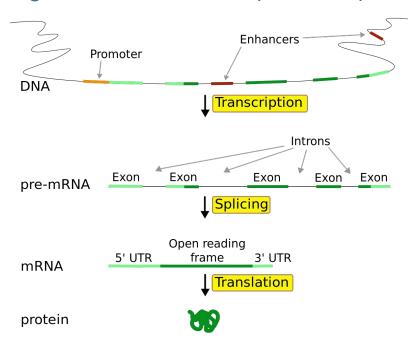
Human genome stats

- The human genome has 23 pairs of chromosomes:
 - Males have 22 autosomes, one X and one Y chromosome
 - Females have 22 autosomes and two X chromosomes
- There is also a separate Mitochondrial DNA contig
- Total length is ~3 Gigabases (3B basepairs)
- About 20 000 genes covering about 30 Megabases
 - ■2% is the coding region



Human exome stats

- Exome part of the human genome coding for proteins/mRNA
- Covers only 2% of genome much cheaper to sequence



Human genome HG38

- Current version of the chromosomes is HG38 (though 37 is still used sometimes)
- Released December 24th 2013
- Additional sequences are added to the genome:
 - Unplaced sequences (some genomes contain them, somewhere)
 - Unlocalized sequences (chromosome known, but coordinates are not)
 - Alternate sequences (some genomes contain them, instead of some parts)
 - Human Herpesvirus 4 type 1
- Patches are commonly added
- http://www.ncbi.nlm.nih.gov/projects/genome/assembly/grc/human/

FASTA file format

- A simple format for storing reference files
- Two "types" of lines:
 - Description lines, start with '>', contain sequence name and optional description
 - Sequence lines follow a description line and contain the actual nucleotide sequence
- Sequences are represented by a single description line,
 followed by one or more sequence lines (usually 70 bases or less in a line)

>chr 1

FASTA File IUPAC Codes (and contig names)

- A, C, T, G, U Nucleotides (Adenine, Cytosine, Guanine, Thymine, Uracil)
- Ambiguous bases:
 - R A or G
 - Y C, T, or U
 - N Any base
 - https://en.wikipedia.org/wiki/FASTA_format
- Often two versions of the Human Reference Genome are found
 - Chromosomes labeled 1, 2, 3... (Human Genome Consortium style)
 - Chromosomes labeled chr1, chr2, chr3... (UCSC style)

FASTA index (fai)

- Some tools build or require indices of the fasta file
 - Needs to be in the same folder, and have the same name as the fasta + .fai extension
- Fai file structure:
 - 1. The name of the sequence
 - 2. The length of the sequence
 - 3. The byte offset of the first base in the file
 - 4. The number of bases in each fasta line
 - 5. The number of bytes in each fasta line

Example:

1	249250621	52	60	61
2	243199373	253404903	60	61
3	198022430	500657651	60	61

Pysam - Python fasta interface

- Pysam a Python toolkit for working with genomic files
- pysam.Fastafile:
 - Create a fasta file parser: fasta = pysam.Fastafile(path_to_file)
 - Get the sequence names in the file: fasta.references
 - Get the lengths of the sequences: fasta.lengths
 - Retrieve a (part of) sequence: fasta.fetch(sequence_name, [start], [stop])
- Fasta coordinates in pysam are zero-based
 - Not true for all file types in that pysam supports
- Other fasta interfaces for Python exist
 - pyfasta works only on fasta files
 - biopython a much larger toolkit that complement sequences, etc.
 - We are describing pysam, as it covers all the file types covered in the course

Exercise: FASTA file format (10 minutes)

- An example FASTA file is found under /sbgenomics/project-files/example_human_reference.fasta
- View the contents of the file
 - You can use "!head filename.ext" in the Notebook to invoke linux head
- Create a pysam Fastafile parser
- Get and print sequence names
 - How many sequences are there?
- Fetch the entire sequence
 - How long is it?
 - Print the first 100 bases

Exercise: FASTA file format (20 minutes)

- Full hg38 FASTA file is located under in /sbgenomics/project-files/Homo_sapiens_assembly38.fasta
- Create a pysam Fastafile parser
- Get sequence names for all contigs
 - How many contigs are there?
 - Read the names of the contigs
 - How long is chromosome 5?
- Fetch section chromosome 17:43044295-43125370
 - What is the "GC content" on this region? (percent of G and C bases)
 - What is the most common 3-mer?
- Fetch base at chromosome 1:248755121
 - What is the base?
- Fetch region at chromosome 1:50000-50100
 - What is the base composition?

FASTQ file format

- Most common format for storing sequencing reads
- It's spread across four lines. The four lines are:
 - ■"@" followed by a read name
 - ■Nucleotide sequence
 - "+", possibly followed by some info, but ignored by most of the tools
 - ■Quality sequence

Example:

@ERR294379.100739024 HS24_09441:8:2203:17450:94030#42/1 AGGGAGTCCACAGTCCAGACTCCACCAGTTCTGACGAAATGATGAGAGCTCAGA

+

BDDEEF?FGFFFHGFFHHGHGGHCH@GHHHGFAHEGFEHGEFGHCCGGGFEGFGFFDFFHB

Annotations file formats

- Annotations are stored in few (similar) file types
 - BED
 - GTF
 - GFF
 - **...**
- **BED file format** (tab-separated):

CHROM START END NAME score ticks blocks

- Chrom, start and end are required
- Start is 0-based
- End is open interval (not included)